

Statistics Introduction

Christof Angermueller

November 27, 2014

1 Clustering


2 PCA

1 Clustering

2 PCA

Clustering

- **Goal:** Finding groups of related items
- How do define relatedness?
- Which clustering methods exist?
- What is Hierarchical clustering?
- How to visualize clustering?



How to define relatedness?

Distance

- Euclidean distance
- Manhattan distance
- Binary distance

Similarity


- Identity
- Correlation
- → Every similarity can be converted into distance

Euclidean distance

Definition

Euclidean distance

$$d(x, y) = \sqrt{\sum_i (x_i - y_i)^2}$$




Manhattan distance

Definition

Manhattan distance

$$d(x, y) = \sum_i |x_i - y_i|$$



Methods

Partitioning clustering


Density clustering

Hierarchical clustering

Methods

Partitioning clustering

- Partition points into k clusters
- k known a-priori
- k-means
- k-medoids




Density clustering

Hierarchical clustering

Methods


Partitioning clustering

- Partition points into k clusters
- k known a-priori
- k-means
- k-medoids



Density clustering


- Cluster points into dense regions
- k unknown a-priori
- DBSCAN
- OPTICS



Hierarchical clustering


Partitioning clustering

- Partition points into k clusters
- k known a-priori
- k-means
- k-medoids




Density clustering

- Cluster points into dense regions
- k unknown a-priori
- DBSCAN
- OPTICS



Hierarchical clustering

- Find hierarchy of clusters
- k unknown a-priori
- Single-linkage
- Complete-linkage
- Average-linkage



Hierarchical clustering

- Constructs hierarchy of clusters represented by a **Cluster dendrogram**
- **Cluster dendrogram**
 - ▶ Leaf nodes: single data points
 - ▶ Inner nodes: cluster of points
 - ▶ Root node: all data points




Algorithm


- 1 Start with clusters that contain only one point
- 2 Compute distance between clusters
- 3 Merge the two clusters with lowest distance into new cluster
- 4 Go to step 2 until one single cluster remains

How to compute the distance between clusters?


- Simple linkage



- Average linkage




- Complete linkage




Data

```
set.seed(5)
n = 5
nclust = 3
x_means = c(1, 2, 2.5)
y_means = c(1, 2.5, 2)
x = rnorm(n * nclust, rep(x_means, each=n), sd=.4)
y = rnorm(n * nclust, rep(y_means, each=n), sd=.4)
clust = factor(rep(1:nclust, each=n))
id = 1:(n * nclust)
cdata = data.frame(id=id, x=x, y=y, clust=clust)
ggplot(cdata, aes(x=x, y=y)) + geom_point(aes(color=clust), size=3, show_guide=F) +
  geom_text(aes(label=id, color=clust), show_guide=F, vjust=-.3, hjust=0)
```




Data

```
set.seed(5)
n = 5
nclust = 3
x_means = c(1, 2, 2.5)
y_means = c(1, 2.5, 2)
x = rnorm(n * nclust, rep(x_means, each=n), sd=.4)
y = rnorm(n * nclust, rep(y_means, each=n), sd=.4)
clust = factor(rep(1:nclust, each=n))
id = 1:(n * nclust)
cdata = data.frame(id=id, x=x, y=y, clust=clust)
ggplot(cdata, aes(x=x, y=y)) + geom_point(color='black', size=3, show_guide=F) +
  geom_text(aes(label=id, color='black', show_guide=F, vjust=-.3, hjust=0))
```




Complete linkage

```
cdist = dist(cbind(cdata$x, cdata$y), method='euclidean')
chclust = hclust(cdist, method='complete')
plot(chclust, ylab='Distance', xlab=NA, main=NA, sub=NA)
```



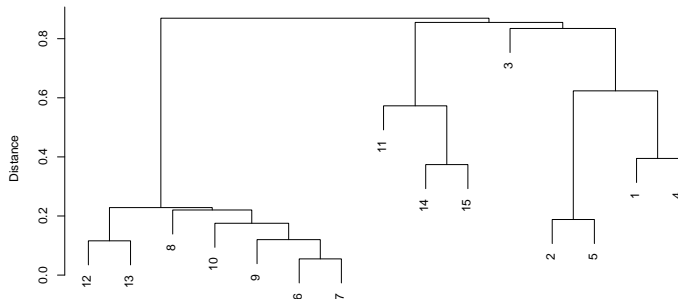
Average linkage

```
cdist = dist(cbind(cdata$x, cdata$y), method='euclidean')  
chclust = hclust(cdist, method='average')  
plot(chclust, ylab='Distance', xlab=NA, main=NA, sub=NA)
```




Single linkage

```
cdist = dist(cbind(cdata$x, cdata$y), method='euclidean')  
chclust = hclust(cdist, method='single')  
plot(chclust, ylab='Distance', xlab=NA, main=NA, sub=NA)
```



Data

```
set.seed(5)
n = 5
nclust = 3
x_means = c(1, 2, 2.5)
y_means = c(1, 2.5, 2)
x = rnorm(n * nclust, rep(x_means, each=n), sd=.4)
y = rnorm(n * nclust, rep(y_means, each=n), sd=.4)
clust = factor(rep(1:nclust, each=n))
id = 1:(n * nclust)
cdata = data.frame(id=id, x=x, y=y, clust=clust)
ggplot(cdata, aes(x=x, y=y)) + geom_point(aes(color=clust), size=3, show_guide=F) +
  geom_text(aes(label=id, color=clust), show_guide=F, vjust=-.3, hjust=0)
```




1 Clustering

2 PCA


High-dimensional data

- Small n large p
- Few samples n (≈ 100)
- Many features p
 - ▶ > 10000 genes
 - ▶ > 270000000 CpG sites



Problems

- High storage costs (memory)
- High computational costs (time)
- Visualization?
- Curse of dimensionality




Principle Component Analysis

- Dimensionality reduction
- Visualization
- Missing values imputation
- Latent factors estimation:
 - ▶ Population structure
 - ▶ Batch-effects
 - ▶ Cell-cycle

Principle components

- Minimize projection error
- Maximize variance




Pearson correlation coefficient

- Measures **linear** dependency between x and y
- $cor(x, y) \in [-1, +1]$
- $cor(x, y) = 0$: no correlation
- $cor(x, y) = -1$: negative correlation
- $cor(x, y) = +1$: positive correlation

Definition (Pearson correlation coefficient)


$$cor(x, y) = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_i (x_i - \bar{x})^2} \sqrt{\sum_i (y_i - \bar{y})^2}}$$

Pearson correlation coefficient




Spearman correlation coefficient

- Measures **monotonic** dependency between x and y
- Pearson correlation coefficient on rank of variables
- $cor(x, y) \in [-1, +1]$



Embryonic Stem Cell (ESC) data

- Single-cell RNA-seq
- $n = 182$ Embryonic stem cells
- $p = 9571$ Genes
- Cell-cycle via Hoechst staining




Genes

Correlation between two ESC

```
c1 <- esc$counts[which(esc$cycle == 'G1')[1],]  
c2 <- esc$counts[which(esc$cycle == 'S')[1],]
```

```
qplot(c1, c2, xlab='G1 cell', ylab='S cell')
```



```
cor(c1, c2,  
    method='pearson')
```


```
## [1] 0.547453
```

```
cor(c1, c2,  
    method='spearman')
```

```
## [1] 0.5785747
```


Correlation matrix

```
cor_cells <- cor(t(esc$counts), method='pearson')  
heatmap(cor_cells, Rowv=NA, Colv=NA)
```



PCA


```
svd <- svd(scale(esc$counts))
dsvd <- as.data.frame(svd$u[, 1:10])
colnames(dsvd) <- paste0('pc', seq(1, ncol(dsvd)))
dsvd$cycle <- esc$cycle
ve <- svd$d^2 / sum(svd$d^2)
barplot(ve[1:10], xlab='Principle components', ylab='% Variance explained',
        col='lightblue')
```



Principle components


PC1 versus PC2

```
ggplot(dsvd, aes(x=pc1, y=pc2, color=cycle)) + geom_point(size=3)
```



Correlation PC1 and cell-cycle

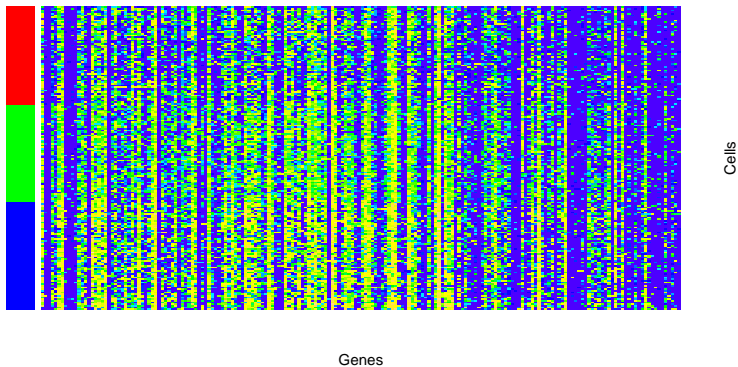
```
ggplot(dsvd, aes(x=cycle, y=pc1, color=cycle)) + geom_boxplot() + geom_jitter(position=position_jitter(width=.1
```




Accounting for cell-cycle

```
svd_compress <- function(svd, k=1) {  
  return (svd$u[, 1:k, drop=FALSE] %*% diag(svd$d[1:k], nrow=k) %*% t(svd$v)[1:k,, drop=FALSE])  
}  
counts_pc1 <- svd_compress(svd, 1)  
counts_npc1 <- esc$counts - counts_pc1
```


Raw counts




Counts explained by cell-cycle





Counts after adjusting for cell-cycle





Further readings

 David J. Balding and Christopher M Bishop.
Handbook of statistical genetics.
Wiley, Chichester [u.a.], 2007.

 Coursera.
Data analysis and statistical inference.

 Coursera.
Mathematical biostatistics boot camp 1.

 Bernard Rosner.
Fundamentals of biostatistics.
Brooks/Cole, Cengage Learning, Boston, 2011.

 Cosma Rohilla Shalizi.
Advanced data analysis from an elementary point of view.
Preprint of book found at [http://www. stat. cmu. edu/ cshalizi/ADAfaEPoV](http://www.stat.cmu.edu/cshalizi/ADAfaEPoV), 2013.

Questions?